

Building Advanced Streaming Pipelines Using Structured Streaming



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Selections, projections and aggregations on streaming data

Adhoc SQL queries on streams

Windowing allows operating on a subset of streaming data

Work with Twitter streaming data

Lateness is the difference between event time and processing time

Watermarking helps deal with lateness

Demo

Using structured streaming in append mode

Demo

Using structured streaming in complete mode

Demo

Aggregations on streaming data

Demo

Running SQL queries on streaming data

Demo

**Including timestamps to mimic event
time**

Demo

Grouping data on timestamps

Stateful Operations on Windows

Transformations



Stateless

Transformations which are applied on a single stream entity



Stateful

Transformations which accumulate across multiple stream entities

Window Transformations



**Accumulate information
across a window in a stream**



85mph

Stateless Transformations




Each entity is operated on
standalone


Speed exceeded? **Alert** triggered



Streaming Data




Streaming Data




Window Transformations

A window is a **subset** of a stream based on

- Time interval
- Count of entities
- Interval between entities



Window Transformations



Transformations can be applied on all entities **within** a window

- sum, min, max, average

Types of Windows

Types of Windows

Tumbling Window

Sliding Window

Count Window

Session Window

Global Window

Types of Windows

Tumbling Window

Sliding Window

Count Window

Session Window


Global Window

Types of Windows



A stream of data

Tumbling Window




Fixed window size

Non-overlapping time


**Number of entities differ within
a window**

Tumbling Window




The window tumbles over the data, in a non-overlapping manner

Tumbling Window



Apply the `sum()` operation on each window

Sliding Window




Fixed window size

Overlapping time - sliding interval

Number of entities differ within a window

Sliding Window

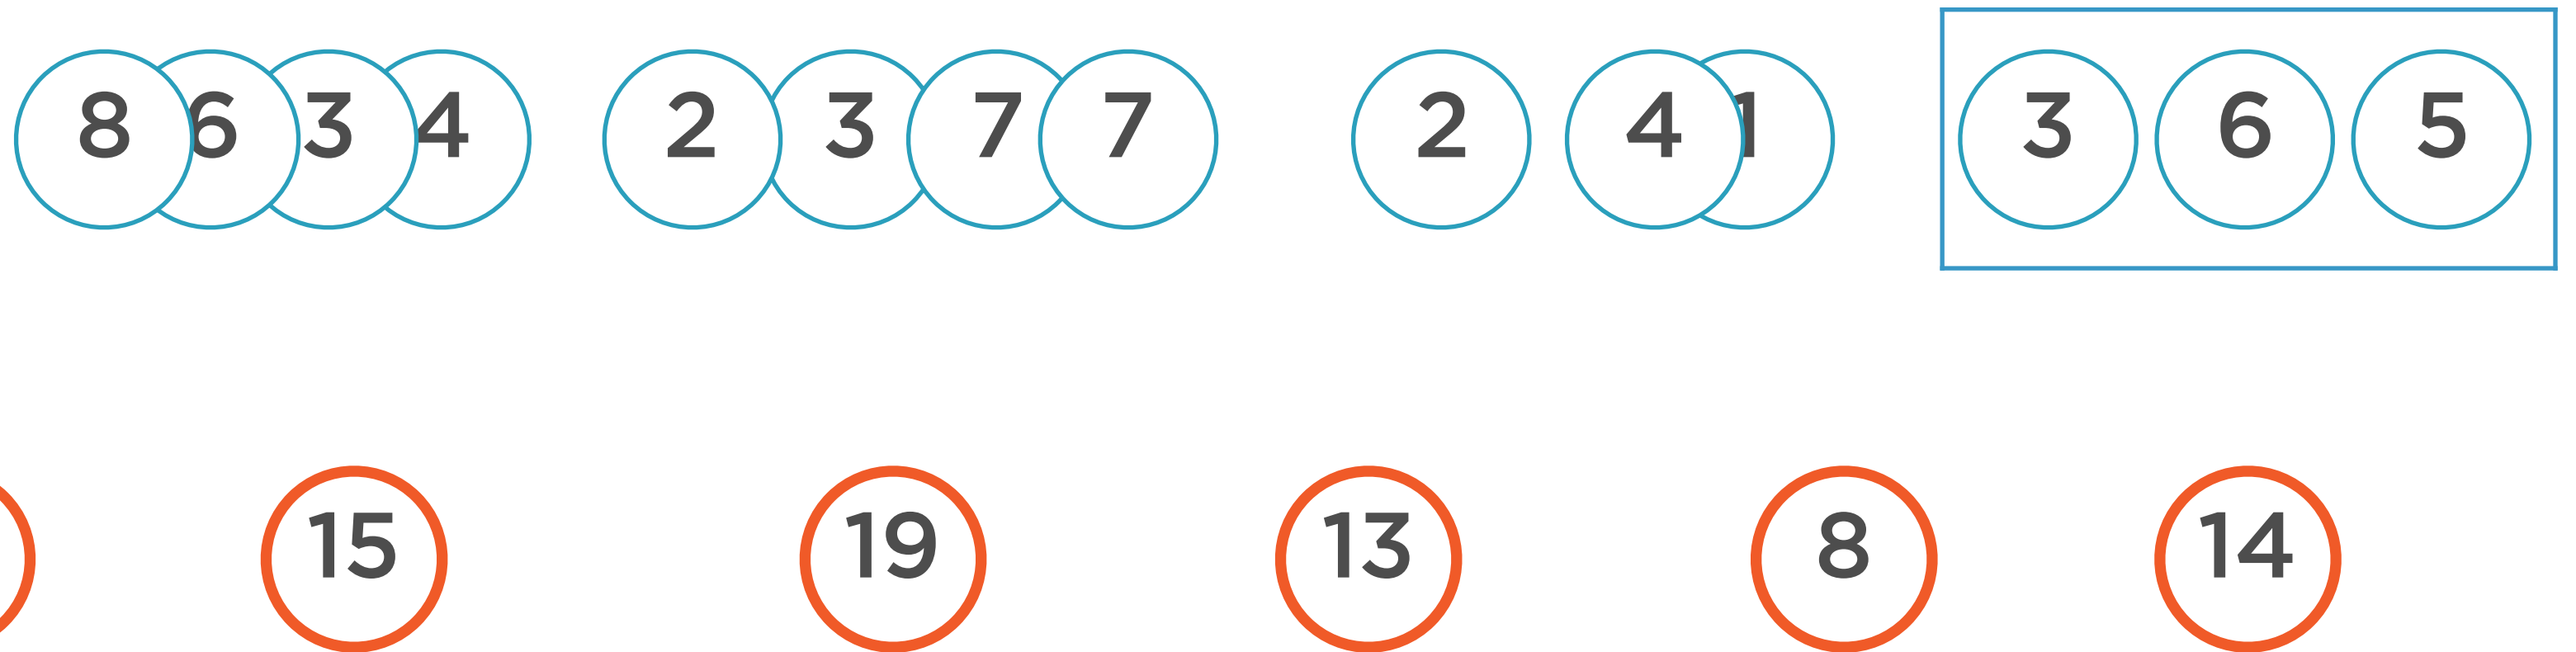


Fixed window size

Overlapping time - sliding interval


Number of entities differ within a window

Sliding Window



Apply the `sum()` operation on each window

Global Window




All data in the stream in one window

The Notion of Time

Time-based Windows

Tumbling and sliding windows consider entities in a fixed interval of **time**



Time-based Windows

Tumbling and sliding windows consider entities in a fixed interval of **time**

There are **different notions of time** that can apply to entities in a stream


Time



Event Time



Ingestion Time



Processing Time



Event Time

The time at which the **event occurred** at its **original source**

- Mobile phone, sensor, website

Usually **embedded within** records

Gives correct results in case of out of order or late events

Ingestion Time



The time at which the event **enters** Spark via a source

Timestamp given by system
chronologically after the event time

Cannot handle out of order events



Processing Time

The **system** time of the machine
processing entities

Chronologically after event time and
ingestion time

Non-deterministic, depends on when
data arrives, how long operations take

Simple, no coordination between
streams and processors


Time



Event Time



Ingestion Time




Processing Time


Time




5



7



Time



Time



Time



Time



Window operations in structured
streaming use **event time**

Demo

Using sliding windows on streaming data

Watermarks and Lateness

How Late is Late?



Class At 9 am

Class starts when
clock strikes 9



Is 9:01 Late?

Realistically, at least
some folks are going
to be a minute late



Is 10:10 late?

A student is an hour
late - allow in or send
back?

Allowed Lateness



The professor “knows” what lateness is reasonable

Students entering within this reasonable lateness are late but OK

Students entering after this reasonable lateness are too late

Excessive Lateness



A student is too late

- Option 1: Send back home
- Option 2: Allow in, continue class
- Option 3: Allow in, restart class(!)

How Late is Late?



Trigger

Class starts when
clock strikes 9



Allowed Lateness

Realistically, at least
some folks are going
to be a minute late



Unacceptable Lateness

A student is an hour
late - allow in or send
back?

Allowed Lateness



The system “knows” what lateness is reasonable

Data entering within this reasonable lateness is late but OK

Data entering after this reasonable lateness is too late

Watermarks and Late Data

The system
“knows” what
lateness is
reasonable

Data entering
within this
reasonable
lateness is late but
OK

Data entering after
this reasonable
lateness is too late

Watermarks and Late Data

Watermark

Threshold of allowed
lateness (event time)

Late Data

Data within watermark is
aggregated (used in
processing Result Table)

Dropped Data

Data outside watermark is
dropped (not used in
processing Result Table)

```
windowedCounts = words.groupBy(  
    window(words.timestamp, "10 minutes", "5 minutes"),  
    words.word  
) .count()
```

Simple Group-by Without Watermark

Count words in each sliding window of width 10 minutes, sliding by 5 minutes

```
windowedCounts = words \  
    .withWatermark("timestamp", "12 minutes") \  
    .groupBy(  
        window(words.timestamp, "10 minutes", "5 minutes"),  
        words.word) \  
    .count()
```

Simple Group-by With Watermark

We define the watermark i.e. lateness threshold to be 12 minutes

```
windowedCounts = words \  
    .withWatermark("timestamp", "12 minutes") \  
    .groupBy(  
        window(words.timestamp, "10 minutes", "5 minutes"),  
        words.word) \  
    .count()
```

Simple Group-by With Watermark

Now window triggering will be delayed by 12 minutes



Watermark

System generated or user specified

If, say network speed drops, watermark can become more lenient

$\text{Lateness} = \text{Processing Time} - \text{Event time}$



Output Modes

Append Mode: Window not triggered at all until watermark elapses

- No partial updates

Update mode: Window will trigger even before watermark elapses

- Engine will keep partial counts

Complete mode: Can not be used with watermarks



Restrictions

No complete-mode queries

Aggregation must be event-time, or event-time window

`.withWatermark` must be called on same timestamp column as aggregate

`.withWatermark` must be called before the aggregation

One-way Guarantee



All data before watermark will definitely not be dropped

All data after watermark may or may not be dropped

Demo

Setting up a Twitter account

**Getting keys and access tokens to
access the Twitter streaming API**

Demo

**Using Tweepy to work with Twitter
streaming data**

Demo

**Count hashtags on Twitter to find
overall trends**

Demo

Count hashtags to find trends using windows

Demo

Count hashtags to find trends using windows

Demo

Join operations using batch and streaming data

Demo

Join operations using aggregations

Demo

Find aggregate ratings using joins

Demo

Join operations on windowed streams

Summary

Selections, projections and aggregations on streaming data

Adhoc SQL queries on streams

Windowing allows operating on a subset of streaming data

Lateness is the difference between event time and processing time

Watermarking helps deal with lateness